

Instruction Fine-tuning using Alpaca Best Practices and Learned Lessons

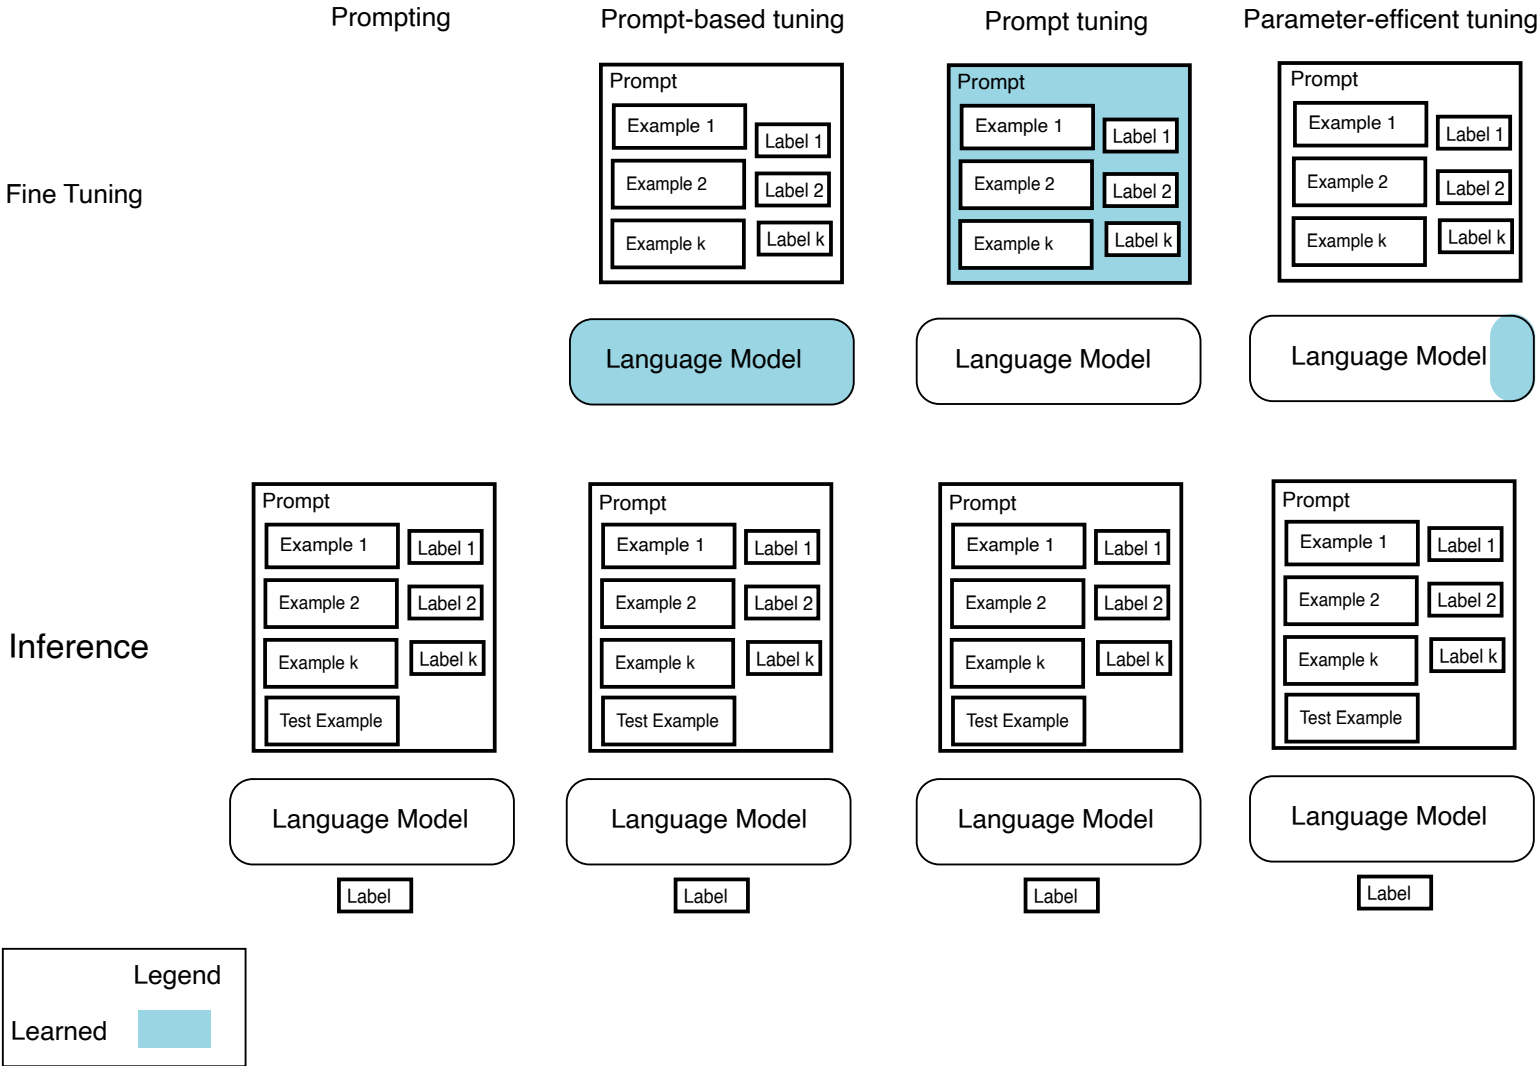
Motivation

Large language models are
few shot and multitask learners

Outlines


- ❑ Prompt-based learning methods
- ❑ Instruction fine-tuning
- ❑ LoRa: Low-Rank Adaptation of Large Language Models
- ❑ Use Case: Instruction fine-tuning on Sentiment Analysis using Alpaca
- ❑ Lessons Learned

Prompt-based learning




Instruction Fine-tuning

Prompt fine-tune



Instruction Fine-tuning

Prompting




175 Tasks

GPT3

Generate



Prompt fine-tune

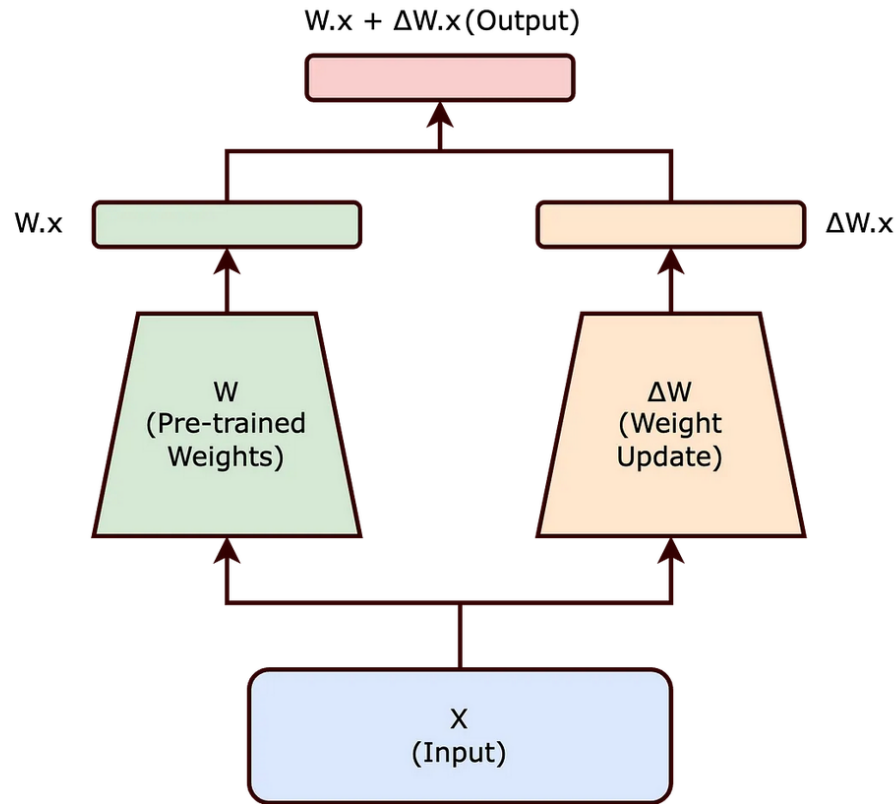


52k Tasks

Alpaca 7B

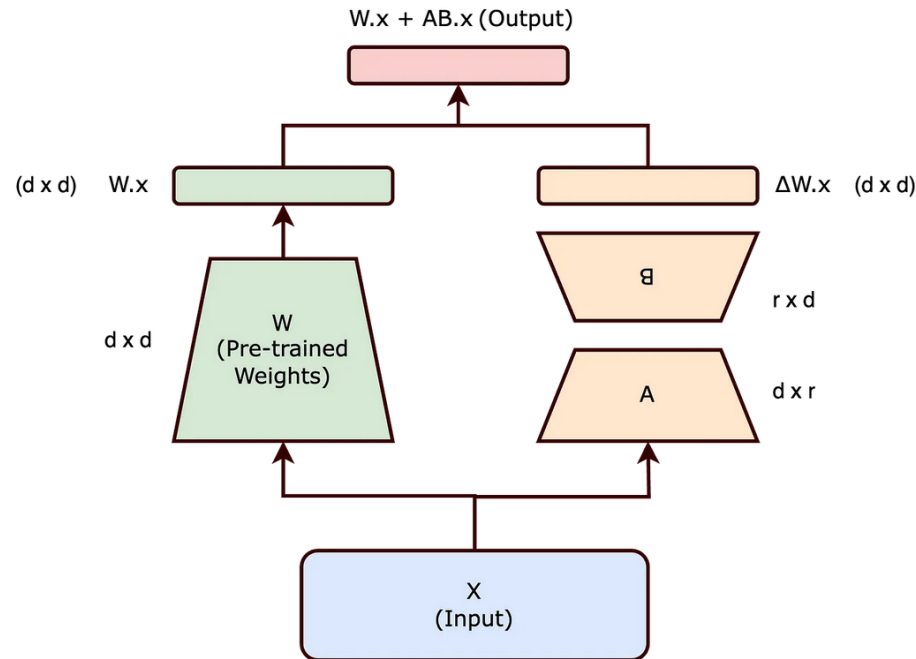
LoRa: Low-Rank Adaptation of Large Language Models

$$h = W_0x + \Delta Wx$$



LoRa: Low-Rank Adaptation of Large Language Models

$$h = W_0x + BAx$$



Implementation

- ❑ Model: wxjiao/alpaca-7b
- ❑ Github: <https://github.com/tloen/alpaca-lora/>

Hardware

- ❑ Nvidia A100 with 80GB
- ❑ 8-cores CPU with access to 128 GB RAM.

Details

- ❑ Float16 precision
- ❑ Early stopping on the training loss
- ❑ Linear Learning rate scheduler
- ❑ Adam Optimizer

Sentiment Analysis

Develop a classifier to categorize the sentiment of a review into positive or negative

Experimental Setup

Amazon Polarity Dataset

split	Reviews	Positive	Negative
Training	32,000	16,022	15,978
Validation	100	54	46
Test	4,000	2,007	1,993

Data Sample

Review 1: I have been programming in C++ for two years on the college level and found this book to be good at building a foundation knowledge of the language that can be expanded on with other book. I am on day 15 and already have a good start on a business program.

Sentiment: Positive

Prompting

Instruction

Classify the following review into two categories: 1) positive, and 2) negative based on its content, given the following examples.

Examples

Example 1

Example 2

Input

I have been programming in C++ for two years on the college level and found this book to be good at building a foundation knowledge of the language that can be expanded on with other book. I am on day 15 and already have a good start on a business program.

Response

Positive

Instruction fine-tuning

Instruction

Classify the following review into two categories: 1) positive, and 2) negative based on its content.

Input

I have been programming in C++ for two years on the college level and found this book to be good at building a foundation knowledge of the language that can be expanded on with other book. I am on day 15 and already have a good start on a business program.

Response

Positive

Sentiment Classification Results with Alpaca-7b

Approach	Shots	Accuracy
SVM		0.83
prompting	16	0.88
Instruction-fine-tuning	16	0.88

Lessons Learned

- ❑ Instruction fine tuning is faster than prompting
- ❑ Set multiple seeds and average over them
- ❑ Use the Alpaca template format for prompt engineering
- ❑ Balance the label distribution in the few shot as well
- ❑ Use sampling methods to choose your few shots
- ❑ Explicitly state the labels in the prompt and define them